

The Little Cloud Foundry Book



By Jonathan ``Duke" Leto

About This Book

License

The Little Cloud Foundry Book book is licensed under the Attribution-NonCommercial-Share-Alike 3.0 Unported license. **You should not have paid for this book.** Donations are totally fine, though :)

You are basically free to copy, distribute, modify or display the book. However, you must attribute the book to Jonathan ``Duke" Leto and Leto Labs LLC and not use it for commercial purposes.

To be clear: if you want to read this book to make you better at doing your job, wonderful. If you would like to use the content of this book in your training program, etc., let's talk. Contact duke@leto.net

You can see the full text of the license at:

<http://creativecommons.org/licenses/by-nc-sa/3.0/legalcode>

About The Author

Jonathan ``Duke" Leto, Founder Leto Labs LLC

Thanks, Y'all

...

Latest Version

The latest source of this book is available at:

<http://github.com/letolabs/the-little-cloud-foundry-book>.

Who Is This Book For?

This book is for three sets of people:

- Application developers wanting to run code *on* Cloud Foundry
- Dev-Ops people who want to install, maintain and manage Cloud Foundry instances
- Developers wanting to contribute, fix bugs and hack on Cloud Foundry itself

If you are an application developer, start from the beginning of this book. If you are a dev-ops person who is already familiar with running applications on CF, you can skip to the ``Installation" chapter.

Developers who are familiar with Cloud Foundry but want to know how to contribute can skip to ``Hacking on Cloud Foundry".

Introduction

Graecum est; non legitur -- ``It's Greek to me"

What is a Platform-As-A-Service (PaaS)?

A PaaS is middleware between the low-level infrastructure (i.e. the actual CPUs, hard-drives, RAM, etc) and the high-level applications that run on top of the platform. The boundary between the platform and the application is clear, but the boundary between the infrastructure and the platform is still largely in flux.

```
-----  
Software-as-a-service  
-----  
Platform-as-a-service  
-----  
Infrastructure-as-a-service  
-----
```

While there are many different definitions of PaaS floating around the market currently, there are a few functions that are specifically the task of PaaS: * live deploying web application code * allocating server resources dynamically * server monitoring * sending http requests to IaaS servers (this needs work)

What is Cloud Foundry?

The name ``Cloud Foundry" can be used to refer to either the still beta commercial VMware product at cloudfoundry.com or the Cloud Foundry Open Source Project. In this book, ``Cloud Foundry" will always refer to the open source project, unless otherwise stated.

Cloud Foundry consists of a large amount of Ruby code along with a small amount of C and Shell scripts that run on top of Ubuntu LTS. The C code is used to manage Linux containers at the kernel level using a similar strategy as the Linux kernel user space program `lxc`. CF also consists of many configuration files in YAML format.

Features

Supported Languages

Supported languages in mainline:

- Ruby (MRI 1.8.x, MRI 1.9.x, JRuby)
- Java
- NodeJS 0.4.x, 0.6.x
- PHP 5.3, 5.4 (ActiveState, AppFog, paas.io)
- Erlang

Supported languages in forks:

- Perl 5 (ActiveState)
- .NET (IronFoundry)
- .NET (Uhuru)
- Haskell 2011.4.0 (paas.io)

Supported Frameworks

Mainline:

- Rails
- Sinatra
- Rack (contributed recently by paas.io)
- Java Spring

Forks:

- Catalyst (ActiveState)
- Dancer (ActiveState)

Glossary

ACM

Access Control Manager. A general system for implementing access control features for applications.

caldecott

A Ruby gem which allows you to ``tunnel" into the various services attached to your application. As an example, you can use caldecott to connect to the mysql command-line of your application running in the cloud with a command similar to

```
vmc tunnel mysql-12345
```

cloud controller

The Cloud Controller can be thought of as the maestro or orchestrator. It also includes the endpoints which vmc communicates with.

Chef

Chef is an open-source systems integration framework to automate cloud-related tasks.

DEA

Droplet Execution Agent. A process which manages the running application instances.

dev-ops

The boundary layer between development and operations is referred to as ``dev-ops". It is characterized by the best practices of operations and development coming together, TODO: EXAMPLE.

droplet

An application, along with all dependencies, in a compressed archive file.

Gerrit

A code review tool based on Git. Gerrit was originally written for the Android Open Source Project at Google and allows many kinds of complex requirements and business rules to be taken into account when reviewing code.

IAAS

Infrastructure-as-a-service

health

Applications can exhibit different states of ``health" which include:

- healthy
- stopped
- flapping
- restarting

health manager

The subsystem of CF which monitors the health of applications and classifies them into a certain number of ``states'', such as ``healthy" or ``stopped".

IAAS

Infrastructure-as-a-service. This is the layer below Platform-as-a-service. An example is Amazon EC2, which takes care of the ``infrastructure" of the actual physical hardware.

lxc

Linux ``containers" which are similar to FreeBSD ``jails". They allow many ``virtual" instances of Linux to run inside of a ``host" Linux, which is actually running on hardware. These have been available in the mainline Linux kernel since 2.6.29.

manifest

A Manifest is a YAML configuration file which lists various properties and metadata about the application, such as endpoint URLs, package versions and checksums.

PAAS

Platform-as-a-service. Examples of this are Cloud Foundry, OpenShift, Rackspace, Google App Engine and many more.

router

The Router takes HTTP requests and sends them to the appropriate running application instance. It usually sits behind one or more load balancers.

stager

The stager takes an application, which is often a directory of files on disk, and turns it into a package that can be deployed to an application instance as part of an update or initial deploy.

Ubuntu LTS

Long-term supported releases of Ubuntu are supported for 2 years by Canonical.

UAA

User Account and Authentication. UAA uses OpenID Connect for authentication (also known as Single Sign On) and the OAuth2 protocol for granting access to resources.

vcap

VMware's Cloud Application Platform. This is the central Git repository which contains the Cloud Foundry codebase.

vmc

Stands for ``VMware Cloud" or ``VMware controller" or whatever you want it to mean, really. It is a ruby gem which is the command-line client to endpoints which implement the Cloud Foundry API.

warden

Warden is the security subsystem. It manages Linux containers, including their creation, destruction and monitoring. It can be thought of as a delicious layer of Ruby on top of a bit of C which uses Linux kernel hooks.

Community

Many individuals and companies have contributed to Cloud Foundry to make it what it is today. Things that were added by the community include PHP, Rack, etc...

In the past, VMware was overwhelmed with Github pull requests, so some of them went seemingly ignored. This is explained by the fact that it was VMware's practice to sync from their private Gerrit repos to their public Github mirror roughly monthly which greatly increases the likelihood of merge conflicts.

To the delight of many, a public Gerrit instance was recently announced. This will allow internal VMware CF developers and external CF developers to work together, in public, which is a huge step in the right direction for VMware. <http://reviews.cloudfoundry.org> Developers can sign in with OpenID or a Google Account and participate in the development of Cloud Foundry.

There is an ``unofficial" cloudfoundry IRC channel on Freenode at #cloudfoundry.

Deploying your App to Cloud Foundry

If you just want to deploy apps to a Cloud Foundry instance, you just need to install a single Ruby gem called ``vmc"

```
gem install vmc
```

Installing Cloud Foundry

Simplest

ActiveState Stackato provides the easiest way to get a working Cloud Foundry instance running locally. There is a single, simple, beautiful command which takes care of everything, including creating a pristine VM:

```
curl get.stackato.com/microcloud | bash
```

Cloud Foundry Open Source Code Flow




Figure 1: CF Community Process

From Github

On a 64bit Ubuntu LTS (10.04.2 works well) with at least 1GB RAM

```
sudo apt-get install openssh-server curl  
bash < <(curl -s -k -B http://git.io/vcap_dev_setup)
```

History of Cloud Foundry

Case Studies

ql.io

node.js grumbles

CF multi-node using AWS

Load Balancer: AWS elastic load balancer

Routers: 2 small instances (1.7 GB of RAM each)

DEA: 2 2xlarge instances (32 GB of RAM each), 1 xlarge instance (15 GB of RAM)

CC/HM/Nats: 1 xlarge instance (15 GB of RAM)

Database: AWS relational database service

NTT

Contributed memcached pull request on Github, has large internal CF cloud.

Links

<http://cloudfoundry.org>

<https://github.com/cloudfoundry>

<http://apidocs.cloudfoundry.com>